

Quantitative Text Analysis and Natural Language Processing using Python

Day 2

Joshua Cova and Luuk Schmitz

2026-01-23

Overview

Part 1: The building blocks of NLP

Part 2: Classification metrics

Coding Exercise 1: Building a DTM

Part 3: Frequency analysis in practice

Part 4: The limits of counting

Coding Exercise 2: Document similarity

Part 5: Setting the scene for word embeddings & LLMs

Part 1: The building blocks of NLP: From text to matrix

The fundamental challenge

We want to use statistical techniques on text to measure the frequency and valence of concepts.

But text is:

- Variable length
- Unstructured
- Symbolic (words, not numbers)

We need a **representation** that converts text to numbers.

Converting text into units

Raw text

- Preprocessing (tokenization, lemmatization, stopwords)
- Representation (DTM, TF-IDF)
- Analysis (frequency, similarity, classification)

Computers work by processing texts into numbers.

The first step therefore has to be to convert raw text into a format that computers can work with

→ **Tokenization**


Structured representation of texts, break the text up into smaller units that a processor can compute

Tokenization example

Large Language Models (LLMs), such as GPT-3 and GPT-4, utilize a process called tokenization. Tokenization involves breaking down text into smaller units, known as tokens, which the model can process and understand. These tokens can range from individual characters to entire words or even larger chunks, depending on the model. For GPT-3 and GPT-4, a Byte Pair Encoding (BPE) tokenizer is used. BPE is a subword tokenization technique that allows the model to dynamically build a vocabulary during training, efficiently representing common words and word fragments. Although the core tokenization process remains similar across different versions of these models, the specific implementation can vary based on the model's architecture and training objectives.

Tokenization example. Source: [Medium](#)

Different types of units



Corpus: All electoral manifestos in Germany in the 2025 federal election

Document: The electoral manifesto of the SPD in the 2025 German federal election

Sentence: “Der Mangel an Fachkräften darf nicht zur Achillesferse unseres Bildungssystems werden.”

Word: “Mangel”

Token: “m”; “angel”

Pre-processing


Common steps:

- Lowercasing (“Democracy” vs. “democracy”)
- Removing punctuation (“election!” vs. “election”)
- Removing numbers
- Removing non-words (signs, html tags)

Goal:

- Reduce noise and keep meaningful variation
- But, as always, choices should depend on context and research question !

Zipf's Law



Zipf's law: Empirical law stating that the frequency of each word is inversely proportional to its rank. Source: Wikipedia

Stop words

- Stop words: Very common words with little substantive meaning (e.g. “the”, “of”)
- Dominate frequency counts
- Several stop words list are available in different Python libraries, but it is advisable to check whether these stop words also make sense for your specific application

Stemming & Lemmatization

Reducing words to their base form

- **Stemming:**
 - Fast
 - Crude
 - “running” → “runni”
- **Lemmatization:**
 - Slower
 - Linguistically informed
 - Better for social science
- **Examples:**
 - “runs”, “running” → “run”
 - “better” → “good”

And now let's see how all of this works in practice...

```
1     # Importing the requests module
2     import requests
3
4     url = "https://www.example.com" # URL of the website
5
6     response = requests.get(url) # Getting the website
7
8     if response.status_code != 200:
9         print(f"Status: {response.status_code} - Try rerunning the code")
10    else:
11        print(f"Status: {response.status_code}\n")
12
13    # Using BeautifulSoup to parse the response object
14    soup = BeautifulSoup(response.content, "html.parser")
15
16    # Finding Post images in the soup
17    images = soup.find_all("img", attrs={"alt": "Post image"})
18
19    # Downloading images
20
```

Part 2: Classification metrics

Connecting to yesterday

Yesterday we discussed **reliability**—do coders agree?

Today we formalize this with **classification metrics**:

- How do we measure agreement between human and machine?
- How do we measure classifier performance?
- What trade-offs do different metrics capture?

The confusion matrix

	Predicted Positive	Predicted Negative
Actually Positive	True Positive (TP)	False Negative (FN)
Actually Negative	False Positive (FP)	True Negative (TN)

Like so much in the social sciences, everything flows from a 2×2 table.

Accuracy: The obvious metric

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

The proportion of all predictions that are correct.

Problem: Can be misleading with imbalanced classes.

Problem 2: Only works as a model-level statistic.

The accuracy trap: An example

Detecting “crisis” language in 1,000 financial reports:

- 50 reports contain crisis language (5%)
- 950 reports don’t (95%)

Classifier A: Always predicts “no crisis” → 95% accuracy!

Classifier B: Correctly identifies 40/50 crises, 10 false alarms → 98% accuracy

Both look good. Only one is useful.

Precision: When false positives are costly

$$\text{Precision} = \frac{TP}{TP + FP}$$

Of all the items we labeled positive, what fraction are actually positive?

High precision = few false alarms

Recall: When false negatives are costly

$$\text{Recall} = \frac{TP}{TP + FN}$$

Of all the items that are actually positive, what fraction did we find?

High recall = few missed cases

Also called **sensitivity** or **true positive rate**.

The precision-recall tradeoff

You usually can't maximize both.

Aggressive classifier: Predicts “positive” liberally

- ↑ Recall (catches more positives)
- ↓ Precision (more false alarms)

Conservative classifier: Predicts “positive” cautiously

- ↓ Recall (misses some positives)
- ↑ Precision (fewer false alarms)

F1 score: Balancing the tradeoff

$$F_1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

The harmonic mean of precision and recall.

- Balances both concerns
- Punishes extreme imbalance
- Widely used as a single summary metric

Worked example

	Predicted Crisis	Predicted Normal
Actually Crisis	40	10
Actually Normal	20	930

Calculate:

- Precision = ?
- Recall = ?
- F1 = ?

Let's run the numbers!

Worked example: Solutions

- **Precision** = $40 / (40 + 20) = 0.67$
 - Two-thirds of predicted crises are real
- **Recall** = $40 / (40 + 10) = 0.80$
 - We catch 80% of actual crises
- **F1** = $2 \times (0.67 \times 0.80) / (0.67 + 0.80) = 0.73$

Is this good? It depends on your research question and your field of study (the stakes in the social sciences \neq clinical trial or astrophysics).

Beyond binary: Multi-class metrics

For multiple categories, calculate per-class then average:

Macro-average: Simple mean across classes

- Treats rare classes equally

Micro-average: Weighted by class frequency

- Treats all instances equally

Weighted average: Weighted by class size

- Common in imbalanced settings

Connecting to inter-coder reliability

Cohen's Kappa for classification:

$$\kappa = \frac{p_o - p_e}{1 - p_e}$$

Same logic as for human coders:

- p_o = proportion where classifier and human agree
- p_e = expected agreement by chance

Use this when comparing automated coding to human gold standard.



Break (15 minutes)

The bag-of-words assumption


Core idea: A document is an unordered collection of words.

Throw away:

- Word order
- Grammar
- Sentence structure

Keep only:

- Which words appear
- How often



What we lose

“The dog bit the man”

↓

{the: 2, dog: 1, bit: 1, man: 1}

“The man bit the dog”

↓

{the: 2, dog: 1, bit: 1, man: 1}

Identical representations!

What we gain

Mathematical tractability

- Documents become vectors
- We can compute distances, similarities
- Standard ML algorithms apply

Scalability

- Simple operations
- Efficient computation
- Works on millions of documents

Surprising effectiveness

- Many tasks care more about “what words” than “what order”
- Topics, sentiment, authorship often work fine

Building the Document-Term Matrix

Step 1: Define your vocabulary (all unique words)

Step 2: For each document, count each word

Step 3: Arrange as matrix

	inflation	economy	growth	crisis
Doc 1	3	2	1	0
Doc 2	1	5	4	0
Doc 3	7	2	0	5

Properties of the DTM

High dimensionality

- Vocabulary of 10,000+ words = 10,000+ columns
- Can be millions for large corpora

Sparsity

- Most cells are zero
- Any one document uses tiny fraction of vocabulary

Zipf's law

- Few words are very common
- Most words are very rare

From counts to weights: TF-IDF

Raw counts treat all words equally.

But “the” appearing 10 times \neq “inflation” appearing 10 times.

TF-IDF reweights by informativeness:

$$\text{TF-IDF}(t, d) = \text{TF}(t, d) \times \text{IDF}(t)$$

TF-IDF: The components

Term Frequency (TF): How often does word t appear in document d ?

$$\text{TF}(t, d) = \frac{\text{count}(t, d)}{\text{total words in } d}$$

Inverse Document Frequency (IDF): How rare is word t across all documents?

$$\text{IDF}(t) = \log \frac{N}{\text{documents containing } t}$$

TF-IDF intuition

Word	Doc 1 TF	Corpus DF	IDF	TF-IDF
the	0.08	1000/1000	0	0
economy	0.04	800/1000	0.22	0.009
stagflation	0.02	10/1000	4.6	0.092

“Stagflation” has lower TF but much higher TF-IDF.

It's distinctive.

The screenshot shows a BBC news article. The header includes the BBC logo, the URL "https://www.bbc.com › news › w...", a link "Vertaal deze pagina", and a three-dot menu icon. The main title of the article is "Does 'lodestar' guide us to anti-Trump op-ed author?". Below the title, the date "6 sep 2018" and a brief description are visible: "The word used in the op-ed has notably been used by none other than Vice-President Mike Pence." The background of the screenshot is dark.

Summary: DTM representations

Type	Cell values	Use case
Binary	0 or 1	Word presence only
Count	Raw frequency	Simple frequency analysis
Normalized	Count / doc length	Comparable across lengths
TF-IDF	Weighted by distinctiveness	Classification, similarity

Coding Exercise 1

Exercise 1: Building a Document-Term Matrix

Goal: Create DTM and TF-IDF representations of the Brexit corpus

Tasks:

1. Load the Brexit speeches data
2. Create a count-based DTM using `CountVectorizer`
3. Create a TF-IDF matrix using `TfidfVectorizer`
4. Examine vocabulary size and sparsity
5. Compare raw counts vs. TF-IDF for specific words
6. Identify most distinctive words per party

Time: 30 minutes, work in pairs



Break (15 minutes)

Part 3: Frequency analysis in practice

From matrix to analysis

Now we have documents as vectors.

What can we do with them?

- **Frequency analysis:** What words are common/rare?
- **Comparison:** How do vocabularies differ across groups?
- **Similarity:** Which documents are alike?
- **Classification:** Can we predict categories?

Word frequency distributions

Zipf's Law revisited: word frequency vs. rank

A few words are very common; most words are rare.

Implications:

- The long tail contains distinctive vocabulary
- High-frequency words often uninformative
- Preprocessing choices matter enormously

Comparing vocabularies across groups

Research question: How do Labour and Conservative MPs talk differently about Brexit?

Approach:

1. Separate corpus by party
2. Calculate word frequencies or TF-IDF per group
3. Identify words that distinguish groups

Keyness: Statistical distinctiveness

- Observed frequency in corpus A
- Expected frequency if A and B used same vocabulary

High keyness = significantly more common in A than B

Commonly used in corpus linguistics (Rayson & Garside 2000)

Limitations of frequency analysis

What frequency captures:

- Word salience (what's talked about)
- Vocabulary differences across groups
- Topic presence

What frequency misses:

- Meaning, context, nuance
- Negation (“not good” = good + not)
- Syntax, relationships between words
- Same word, different meanings (polysemy)

When frequency analysis is enough

- ✓ Measuring visibility/salience of entities or topics
- ✓ Comparing vocabulary across groups
- ✓ Tracking word usage over time
- ✓ Exploratory analysis of corpus content
- ✓ Simple content analysis at scale

Part 4: The limits of counting

The negation problem

Dictionary says “good” is positive.

But:

- “not good” → negative
- “not very good” → negative
- “wouldn’t say it’s good” → negative
- “good for nothing” → negative

Bag of words sees: {good: 1, not: 1, ...}



The polysemy problem

“Bank” means:

- Financial institution (“the bank increased rates”)
- River edge (“sitting by the river bank”)
- Verb for turning (“bank left at the intersection”)
- Reliance (“bank on this strategy”)



Bag of words: bank = bank = bank = bank

The synonymy problem

Different words, same meaning:

- “happy” / “glad” / “joyful” / “pleased”
- “car” / “automobile” / “vehicle”
- “decrease” / “decline” / “drop” / “fall”

Bag of words: all different features

No recognition of semantic similarity

The compositionality problem

Meaning often comes from combination:

- “hot dog” ≠ hot + dog
- “kick the bucket” ≠ kick + bucket
- “white house” ≠ white + house (when capitalized)

Multi-word expressions, idioms, named entities

Revisiting the hot model problem

Despite a hot 2023 and the recent Hansen et al paper, there is still reason to doubt very high climate sensitivity models

What we want from a representation

1. Similar words should be close together

- “good” near “excellent”
- “good” far from “bad”

2. Context should matter

- “bank” in financial context \neq “bank” in nature context

3. Relationships should be preserved

- king:queen :: man:woman

This is what embeddings provide.

Text classification

Simple, but very good use case!

Goal: Use text classification strategies to predict whether an email is ham (0) or spam (1).

Predict classes given the message of a text.

Naive Bayes: $P(\text{Class} \mid \text{Text}) \propto P(\text{Text} \mid \text{Class}) P(\text{Class})$

- Model word probabilities per class
- We choose the class with highest probability
- Assumption that each word contributes independently to the classification is **wrong** but **useful**

Part 5: Setting the scene

Where we've been

Raw text

- Preprocessing (tokenization, lemmatization, stopwords)
- Representation (DTM, TF-IDF)
- Analysis (frequency, similarity, classification)

This pipeline underlies most traditional text analysis.

Limitations: no semantics, no context, no compositionality.

Where we're going

Word embeddings (Word2Vec, GloVe, 2013–)

- Dense vectors that capture semantic relationships
- Pre-trained on massive corpora
- $\text{king} - \text{man} + \text{woman} \approx \text{queen}$

Contextual embeddings (BERT, GPT, 2018–)

- Same word gets different vectors in different contexts
- Pre-trained on next-word prediction
- Fine-tunable for specific tasks

Large Language Models (GPT-4, Claude, 2022–)

- Zero-shot classification through prompting
- In-context learning


Word embeddings: The key insight

Distributional hypothesis: “You shall know a word by the company it keeps” (Firth 1957)

Words appearing in similar contexts have similar meanings.

Train a model to predict context → words with similar contexts get similar vectors.

Word embeddings: Visualization I



Word embeddings: Visualization II

crown
↑
king • queen
|
man • woman
↓
(gender axis)

Semantic relationships become geometric relationships.

$$\text{king} - \text{man} + \text{woman} \approx \text{queen}$$

From words to documents

Simple approach: Average the word vectors

$$\vec{d} = \frac{1}{n} \sum_{i=1}^n \vec{w}_i$$

Better than bag-of-words because word vectors capture similarity.

“Happy” and “joyful” contribute similarly.

Contextual embeddings: The next step

Problem with static embeddings:

“bank” has one vector, regardless of context.

Solution: Make embeddings depend on context.

BERT, GPT, and other transformers learn contextual representations.

“The river bank was muddy” → different “bank” vector

“The investment bank collapsed” → different “bank” vector

LLMs as classifiers

Traditional approach:

Train a classifier on labeled data.

LLM approach:

Just ask the model!

Classify this text as positive, negative, or neutral:

"The fiscal outlook remains challenging despite recent gains."

Answer: Negative

A teaser: Same task, different methods

Brexit speech sentiment:

Method	Approach	Performance
Dictionary (VADER)	Count positive/negative words	Baseline
Bag-of-words + SVM	Train classifier on TF-IDF	Better
Fine-tuned BERT	Transfer learning	Much better
LLM prompting	Zero-shot classification	Comparable or better

The tradeoffs

Method	Interpretability	Data needed	Compute	Performance
Dictionary	High	None	Low	Variable
BoW + ML	Medium	Moderate	Low	Good
Embeddings	Low	Less	Medium	Better
LLMs	Low	None	High	Best*

*For many tasks; highly dependent on prompting.

Practical guidance

Start simple, add complexity only if needed.

1. Establish human baseline (inter-coder reliability)
2. Try dictionary methods if concept is well-defined
3. Try supervised learning if you have labeled data
4. Consider embeddings if little labeled data
5. Use LLMs for exploration and hard cases
6. **Validate everything**

Wrapping up

Key takeaways

1. **Classification metrics** (precision, recall, F1) formalize evaluation
2. **Bag-of-words** is simple but surprisingly effective
3. **TF-IDF** emphasizes distinctive words
4. **Similarity and classification** flow from representation
5. **Embeddings and LLMs** address BoW limitations—coming soon!

What's next: Spring workshop

Machine Learning and LLMs for Text Analysis

- Word embeddings (Word2Vec, GloVe)
- Transfer learning with pre-trained models
- Fine-tuning BERT for classification
- Using LLMs (GPT, Claude) for text analysis
- Prompt engineering for zero-shot classification

Resources for further learning

Books:

- Grimmer, Roberts & Stewart (2022). *Text as Data*
- Jurafsky & Martin (2024). *Speech and Language Processing* (free online)

Papers:

- Gentzkow, Kelly & Taddy (2019). “Text as Data” *JEL*
- Grimmer & Stewart (2013). “Text as Data” *Political Analysis*

Tutorials:

- scikit-learn text tutorial
- spaCy course (free)
- Hugging Face NLP course (free)

Thank you!

Questions?

References

- Firth, J. R. (1957). A synopsis of linguistic theory 1930-55. *Studies in Linguistic Analysis*.
- Gentzkow, M., Kelly, B., & Taddy, M. (2019). Text as data. *Journal of Economic Literature*, 57(3), 535-574.
- Grimmer, J., Roberts, M. E., & Stewart, B. M. (2022). *Text as Data: A New Framework for Machine Learning and the Social Sciences*. Princeton University Press.
- Grimmer, J., & Stewart, B. M. (2013). Text as data: The promise and pitfalls of automatic content analysis. *Political Analysis*, 21(3), 267-297.
- Rayson, P., & Garside, R. (2000). Comparing corpora using frequency profiling. *Proceedings of the Workshop on Comparing Corpora*.